

Assignment is below at the end

- <https://scikit-learn.org/stable/modules/tree.html> (<https://scikit-learn.org/stable/modules/tree.html>)
- <https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html> (<https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>)
- https://scikit-learn.org/stable/modules/generated/sklearn.tree.plot_tree.html (https://scikit-learn.org/stable/modules/generated/sklearn.tree.plot_tree.html)

```
In [164]: import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
plt.rcParams['figure.figsize'] = (20, 6)
plt.rcParams['font.size'] = 14
import pandas as pd
```

```
In [165]: df = pd.read_csv('../data/adult.data', index_col=False)
```

```
In [166]: golden = pd.read_csv('../data/adult.test', index_col=False)
```

```
In [167]: golden.head()
```

Out[167]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	salary
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0	0	40	United-States	<=50K.
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0	0	50	United-States	<=50K.
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	0	0	40	United-States	>50K.
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	7688	0	40	United-States	>50K.
4	18	?	103497	Some-college	10	Never-married	?	Own-child	White	Female	0	0	30	United-States	<=50K.

```
In [168]: df.head()
```

Out[168]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	salary
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K

```
In [169]: df.columns
```

```
Out[169]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',
'marital-status', 'occupation', 'relationship', 'race', 'sex',
'capital-gain', 'capital-loss', 'hours-per-week', 'native-country',
'salary'],
dtype='object')
```

```
In [170]: from sklearn import preprocessing
```

```
In [171]: # Columns we want to transform
transform_columns = ['sex']

#Columns we can't use because non-numerical
non_num_columns = ['workclass', 'education', 'marital-status',
'occupation', 'relationship', 'race', 'sex',
'native-country']
```

First let's try using `pandas.get_dummies()` to transform columns

```
In [172]: dummies = pd.get_dummies(df[transform_columns])
          dummies
```

Out[172]:

	sex_Female	sex_Male
0	0	1
1	0	1
2	0	1
3	0	1
4	1	0
...
32556	1	0
32557	0	1
32558	1	0
32559	0	1
32560	1	0

32561 rows × 2 columns

```
In [173]: dummies.shape
```

Out[173]: (32561, 2)

sklearn has a similar process for OneHot Encoding features

```
In [301]: onehot = preprocessing.OneHotEncoder(handle_unknown="infrequent_if_exist", sparse_output=False)
          onehot.fit(df[transform_columns])
```

Out[301]:

```
OneHotEncoder
OneHotEncoder(handle_unknown='infrequent_if_exist', sparse_output=False)
```

```
In [291]: onehot.categories_
```

Out[291]: [array([' Female', ' Male'], dtype=object)]

```
In [176]: sex = onehot.transform(df[transform_columns])
          sex
```

Out[176]: array([[0., 1.],
[0., 1.],
[0., 1.],
...,
[1., 0.],
[0., 1.],
[1., 0.]])

```
In [177]: sex.shape
```

Out[177]: (32561, 2)

In addition to OneHot encoding there is Ordinal Encoding

```
In [178]: enc = preprocessing.OrdinalEncoder()
          enc.fit(df[["salary"]])
          salary = enc.transform(df[["salary"]])
          salary
```

Out[178]: array([[0.],
[0.],
[0.],
...,
[0.],
[0.],
[1.]])

```
In [179]: enc.categories_[0]
```

```
Out[179]: array([' <=50K', ' >50K'], dtype=object)
```

```
In [180]: x = df.copy()

# transformed = pd.get_dummies(df[transform_columns])

onehot = preprocessing.OneHotEncoder(handle_unknown="infrequent_if_exist", sparse_output=False).fit(df[transform_columns])
enc = preprocessing.OrdinalEncoder()

enc.fit(df[["salary"]])

transformed = onehot.transform(df[transform_columns])
new_cols = list(onehot.categories_[0].flatten())
df_trans = pd.DataFrame(transformed, columns=new_cols)

x = pd.concat(
    [
        x.drop(non_num_columns, axis=1),
        df_trans
    ],
    axis=1,
)

x["salary"] = enc.transform(df[["salary"]])
```

```
In [181]: x.head()
```

```
Out[181]:
```

	age	fnlwtg	education-num	capital-gain	capital-loss	hours-per-week	salary	Female	Male
0	39	77516	13	2174	0	40	0.0	0.0	1.0
1	50	83311	13	0	0	13	0.0	0.0	1.0
2	38	215646	9	0	0	40	0.0	0.0	1.0
3	53	234721	7	0	0	40	0.0	0.0	1.0
4	28	338409	13	0	0	40	0.0	1.0	0.0

```
In [182]: xt = golden.copy()

transformed = onehot.transform(xt[transform_columns])
new_cols = list(onehot.categories_[0].flatten())
df_trans = pd.DataFrame(transformed, columns=new_cols)

xt = pd.concat(
    [
        xt.drop(non_num_columns, axis=1),
        df_trans
    ],
    axis=1,
)

xt["salary"] = enc.fit_transform(golden[["salary"]])
```

```
In [183]: xt.salary.value_counts()
```

```
Out[183]: 0.0    12435
          1.0    3846
          Name: salary, dtype: int64
```

```
In [184]: enc.categories_
```

```
Out[184]: [array([' <=50K.', ' >50K.'], dtype=object)]
```

```
In [185]: from sklearn.tree import DecisionTreeClassifier
          from sklearn.ensemble import RandomForestClassifier
          from sklearn.ensemble import GradientBoostingClassifier
```

Choose the model of your preference: DecisionTree or RandomForest

```
In [186]: model = RandomForestClassifier(criterion='entropy')
```

```
In [187]: model = DecisionTreeClassifier(criterion='entropy', max_depth=None)
```

```
In [188]: model.fit(x.drop(['fnlwgt', 'salary'], axis=1), x.salary)
```

```
Out[188]: DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy')
```

```
In [189]: model.tree_.node_count
```

```
Out[189]: 8313
```

```
In [190]: list(zip(x.drop(['fnlwgt', 'salary'], axis=1).columns, model.feature_importances_))
```

```
Out[190]: [('age', 0.3226394602924674),
 ('education-num', 0.1616023519291502),
 ('capital-gain', 0.22748700351709567),
 ('capital-loss', 0.079214470333594),
 ('hours-per-week', 0.1537256624686072),
 (' Female', 0.0013013658419538756),
 (' Male', 0.05402968561713164)]
```

```
In [191]: list(zip(x.drop(['fnlwgt', 'salary'], axis=1).columns, model.feature_importances_))
```

```
Out[191]: [('age', 0.3226394602924674),
 ('education-num', 0.1616023519291502),
 ('capital-gain', 0.22748700351709567),
 ('capital-loss', 0.079214470333594),
 ('hours-per-week', 0.1537256624686072),
 (' Female', 0.0013013658419538756),
 (' Male', 0.05402968561713164)]
```

```
In [192]: x.drop(['fnlwgt', 'salary'], axis=1).head()
```

```
Out[192]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	Female	Male
0	39	13	2174	0	40	0.0	1.0
1	50	13	0	0	13	0.0	1.0
2	38	9	0	0	40	0.0	1.0
3	53	7	0	0	40	0.0	1.0
4	28	13	0	0	40	1.0	0.0

```
In [193]: set(x.columns) - set(xt.columns)
```

```
Out[193]: set()
```

```
In [194]: list(x.drop('salary', axis=1).columns)
```

```
Out[194]: ['age',
 'fnlwgt',
 'education-num',
 'capital-gain',
 'capital-loss',
 'hours-per-week',
 ' Female',
 ' Male']
```

```
In [195]: list(x)
```

```
Out[195]: ['age',
 'fnlwgt',
 'education-num',
 'capital-gain',
 'capital-loss',
 'hours-per-week',
 'salary',
 ' Female',
 ' Male']
```

```
In [196]: list(xt)
```

```
Out[196]: ['age',
 'fnlwgt',
 'education-num',
 'capital-gain',
 'capital-loss',
 'hours-per-week',
 'salary',
 ' Female',
 ' Male']
```

```
In [197]: predictions = model.predict(xt.drop(['fnlwgt', 'salary'], axis=1))
          predictionsx = model.predict(x.drop(['fnlwgt', 'salary'], axis=1))
```

```
In [198]: from sklearn.metrics import (
          accuracy_score,
          classification_report,
          confusion_matrix, auc, roc_curve
          )
```

```
In [199]: accuracy_score(xt.salary, predictions)
```

```
Out[199]: 0.8205269946563479
```

```
In [200]: accuracy_score(xt.salary, predictions)
```

```
Out[200]: 0.8205269946563479
```

```
In [201]: confusion_matrix(xt.salary, predictions)
```

```
Out[201]: array([[11461,   974],
                 [ 1948,  1898]])
```

```
In [202]: print(classification_report(xt.salary, predictions))
```

	precision	recall	f1-score	support
0.0	0.85	0.92	0.89	12435
1.0	0.66	0.49	0.57	3846
accuracy			0.82	16281
macro avg	0.76	0.71	0.73	16281
weighted avg	0.81	0.82	0.81	16281

```
In [203]: print(classification_report(xt.salary, predictions))
```

	precision	recall	f1-score	support
0.0	0.85	0.92	0.89	12435
1.0	0.66	0.49	0.57	3846
accuracy			0.82	16281
macro avg	0.76	0.71	0.73	16281
weighted avg	0.81	0.82	0.81	16281

```
In [204]: accuracy_score(x.salary, predictionsx)
```

```
Out[204]: 0.8955806025613464
```

```
In [205]: confusion_matrix(x.salary, predictionsx)
```

```
Out[205]: array([[24097,   623],
                 [ 2777,  5064]])
```

```
In [206]: print(classification_report(x.salary, predictionsx))
```

	precision	recall	f1-score	support
0.0	0.90	0.97	0.93	24720
1.0	0.89	0.65	0.75	7841
accuracy			0.90	32561
macro avg	0.89	0.81	0.84	32561
weighted avg	0.90	0.90	0.89	32561

```
In [207]: print(classification_report(x.salary, predictionsx))
```

	precision	recall	f1-score	support
0.0	0.90	0.97	0.93	24720
1.0	0.89	0.65	0.75	7841
accuracy			0.90	32561
macro avg	0.89	0.81	0.84	32561
weighted avg	0.90	0.90	0.89	32561

For the following use the above adult dataset.

1. Show the RandomForest outperforms the DecisionTree for a fixed max_depth by training using the train set and calculate precision, recall, f1, confusion matrix on golden-test set. Start with only numerical features/columns. (age, education-num, capital-gain, capital-loss, hours-per-week)

```
In [212]: x1 = x.copy()
          xt1 = xt.copy()
```

```
In [213]: x1.head()
```

```
Out[213]:
```

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	salary	Female	Male
0	39	77516	13	2174	0	40	0.0	0.0	1.0
1	50	83311	13	0	0	13	0.0	0.0	1.0
2	38	215646	9	0	0	40	0.0	0.0	1.0
3	53	234721	7	0	0	40	0.0	0.0	1.0
4	28	338409	13	0	0	40	0.0	1.0	0.0

```
In [214]: xt1.head()
```

```
Out[214]:
```

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	salary	Female	Male
0	25	226802	7	0	0	40	0.0	0.0	1.0
1	38	89814	9	0	0	50	0.0	0.0	1.0
2	28	336951	12	0	0	40	1.0	0.0	1.0
3	44	160323	10	7688	0	40	1.0	0.0	1.0
4	18	103497	10	0	0	30	0.0	1.0	0.0

```
In [271]: x1_rf1 = RandomForestClassifier(criterion='entropy', max_depth = 7)
          x1_dt1 = DecisionTreeClassifier(criterion='entropy', max_depth = 7)
          x1_rf2 = RandomForestClassifier(criterion='entropy', max_depth = 3)
          x1_dt2 = DecisionTreeClassifier(criterion='entropy', max_depth = 3)
```

```
In [258]: x1_rf1.fit(x1.drop(['salary'],axis=1),x1.salary)
```

```
Out[258]:
```

RandomForestClassifier
RandomForestClassifier(criterion='entropy', max_depth=7)

```
In [259]: x1_dt1.fit(x1.drop(['salary'],axis=1),x1.salary)
```

```
Out[259]:
```

DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=7)

```
In [272]: x1_rf2.fit(x1.drop(['salary'],axis=1),x1.salary)
```

```
Out[272]:
```

RandomForestClassifier
RandomForestClassifier(criterion='entropy', max_depth=3)

```
In [273]: x1_dt2.fit(x1.drop(['salary'],axis=1),x1.salary)
```

```
Out[273]:
```

DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=3)

```
In [260]: list(zip(x1.drop(['salary'], axis=1).columns, x1_rf1.feature_importances_))
```

```
Out[260]: [('age', 0.240093219743165),
          ('fnlwgt', 0.011534989957478658),
          ('education-num', 0.19846941504636445),
          ('capital-gain', 0.26679223958878184),
          ('capital-loss', 0.07051146866383945),
          ('hours-per-week', 0.09498045659518521),
          (' Female', 0.05217555194014917),
          (' Male', 0.06544265846503625)]
```

```
In [261]: list(zip(x1.drop(['salary'], axis=1).columns, x1_dt1.feature_importances_))
```

```
Out[261]: [('age', 0.2829603491199833),
 ('fmlwgt', 0.006306080004978193),
 ('education-num', 0.19704210624159596),
 ('capital-gain', 0.3217591678808323),
 ('capital-loss', 0.05606590281242601),
 ('hours-per-week', 0.042497567987501844),
 (' Female', 0.03501265095704486),
 (' Male', 0.05835617499563742)]
```

```
In [274]: list(zip(x1.drop(['salary'], axis=1).columns, x1_rf2.feature_importances_))
```

```
Out[274]: [('age', 0.24634078131354797),
 ('fmlwgt', 0.0007106486235433378),
 ('education-num', 0.1969002410895295),
 ('capital-gain', 0.28853155826174126),
 ('capital-loss', 0.04514355356598634),
 ('hours-per-week', 0.08644270766813066),
 (' Female', 0.05963320101352708),
 (' Male', 0.07629730846399385)]
```

```
In [275]: list(zip(x1.drop(['salary'], axis=1).columns, x1_dt2.feature_importances_))
```

```
Out[275]: [('age', 0.3423330938633231),
 ('fmlwgt', 0.0),
 ('education-num', 0.2214011618205338),
 ('capital-gain', 0.4357102902795111),
 ('capital-loss', 0.0),
 ('hours-per-week', 0.000555454036632138),
 (' Female', 0.0),
 (' Male', 0.0)]
```

```
In [262]: rf1_pred1 = x1_rf1.predict(xt1.drop(['salary'],axis=1))
```

```
In [263]: dt1_pred1 = x1_dt1.predict(xt1.drop(['salary'],axis=1))
```

```
In [276]: rf2_pred2 = x1_rf2.predict(xt1.drop(['salary'],axis=1))
```

```
In [277]: dt2_pred2 = x1_dt2.predict(xt1.drop(['salary'],axis=1))
```

```
In [264]: accuracy_score(xt1.salary, rf1_pred1)
```

```
Out[264]: 0.8369879000061421
```

```
In [265]: accuracy_score(xt1.salary, dt1_pred1)
```

```
Out[265]: 0.8309686137215159
```

```
In [ ]: accuracy_score(xt1.salary, rf1_pred1)
```

```
In [ ]: accuracy_score(xt1.salary, dt1_pred1)
```

```
In [266]: confusion_matrix(xt1.salary, rf1_pred1)
```

```
Out[266]: array([[12006,   429],
 [ 2225,  1621]])
```

```
In [267]: confusion_matrix(xt1.salary, dt1_pred1)
```

```
Out[267]: array([[11767,   668],
 [ 2084,  1762]])
```

```
In [278]: confusion_matrix(xt1.salary, rf2_pred2)
```

```
Out[278]: array([[12410,    25],
 [ 3032,   814]])
```

```
In [280]: confusion_matrix(xt1.salary, dt2_pred2)
```

```
Out[280]: array([[12428,     7],
 [ 3199,   647]])
```

```
In [279]: print(classification_report(xt1.salary, rf1_pred1))
```

	precision	recall	f1-score	support
0.0	0.84	0.97	0.90	12435
1.0	0.79	0.42	0.55	3846
accuracy			0.84	16281
macro avg	0.82	0.69	0.73	16281
weighted avg	0.83	0.84	0.82	16281

```
In [269]: print(classification_report(xt1.salary, dt1_pred1))
```

	precision	recall	f1-score	support
0.0	0.85	0.95	0.90	12435
1.0	0.73	0.46	0.56	3846
accuracy			0.83	16281
macro avg	0.79	0.70	0.73	16281
weighted avg	0.82	0.83	0.82	16281

```
In [281]: print(classification_report(xt1.salary, rf2_pred2))
```

	precision	recall	f1-score	support
0.0	0.80	1.00	0.89	12435
1.0	0.97	0.21	0.35	3846
accuracy			0.81	16281
macro avg	0.89	0.60	0.62	16281
weighted avg	0.84	0.81	0.76	16281

```
In [282]: print(classification_report(xt1.salary, dt2_pred2))
```

	precision	recall	f1-score	support
0.0	0.80	1.00	0.89	12435
1.0	0.99	0.17	0.29	3846
accuracy			0.80	16281
macro avg	0.89	0.58	0.59	16281
weighted avg	0.84	0.80	0.74	16281

2. Use a RandomForest or DecisionTree and the adult dataset, systematically add new columns, one by one, that are non-numerical but converted using the feature-extraction techniques we learned. Using the golden-test set show [precision, recall, f1, confusion matrix] for each additional feature added.

```
In [296]: x1.head()
```

```
Out[296]:
```

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	salary	Female	Male
0	39	77516	13	2174	0	40	0.0	0.0	1.0
1	50	83311	13	0	0	13	0.0	0.0	1.0
2	38	215646	9	0	0	40	0.0	0.0	1.0
3	53	234721	7	0	0	40	0.0	0.0	1.0
4	28	338409	13	0	0	40	0.0	1.0	0.0

In [297]: `df.head()`

Out[297]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	salary
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K

In [307]: `x2 = x1.copy()`
`x2['marital-status'] = enc.fit_transform(df[['marital-status']])`
`x2.head()`

Out[307]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	salary	Female	Male	marital-status
0	39	77516	13	2174	0	40	0.0	0.0	1.0	4.0
1	50	83311	13	0	0	13	0.0	0.0	1.0	2.0
2	38	215646	9	0	0	40	0.0	0.0	1.0	0.0
3	53	234721	7	0	0	40	0.0	0.0	1.0	2.0
4	28	338409	13	0	0	40	0.0	1.0	0.0	2.0

In [309]: `xt2 = xt1.copy()`
`xt2['marital-status'] = enc.fit_transform(golden[['marital-status']])`
`xt2.head()`

Out[309]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	salary	Female	Male	marital-status
0	25	226802	7	0	0	40	0.0	0.0	1.0	4.0
1	38	89814	9	0	0	50	0.0	0.0	1.0	2.0
2	28	336951	12	0	0	40	1.0	0.0	1.0	2.0
3	44	160323	10	7688	0	40	1.0	0.0	1.0	2.0
4	18	103497	10	0	0	30	0.0	1.0	0.0	4.0

In [311]: `rf3 = RandomForestClassifier(criterion='entropy', max_depth = 3)`
`dt3 = DecisionTreeClassifier(criterion='entropy', max_depth = 3)`

In [314]: `x2_rf = rf3.fit(x2.drop(['salary'], axis=1), x2.salary)`

In [321]: `x2_rf_pred = x2_rf.predict(xt2.drop(['salary'], axis=1))`

In [322]: `print(classification_report(xt2.salary, x2_rf_pred))`

```

              precision    recall  f1-score   support

    0.0         0.83      0.98      0.90      12435
    1.0         0.86      0.33      0.48       3846

 accuracy          0.83      16281
 macro avg          0.84      0.66      0.69      16281
 weighted avg          0.83      0.83      0.80      16281

```

In [318]: `x3 = x2.copy()`
`x3['native-country'] = enc.fit_transform(df[['native-country']])`
`x3.head()`

Out[318]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	salary	Female	Male	marital-status	native-country
0	39	77516	13	2174	0	40	0.0	0.0	1.0	4.0	39.0
1	50	83311	13	0	0	13	0.0	0.0	1.0	2.0	39.0
2	38	215646	9	0	0	40	0.0	0.0	1.0	0.0	39.0
3	53	234721	7	0	0	40	0.0	0.0	1.0	2.0	39.0
4	28	338409	13	0	0	40	0.0	1.0	0.0	2.0	5.0

```
In [325]: xt3 = xt2.copy()
xt3['native-country'] = enc.fit_transform(golden[['native-country']])
xt3.head()
```

Out[325]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	salary	Female	Male	marital-status	native-country
0	25	226802	7	0	0	40	0.0	0.0	1.0	4.0	38.0
1	38	89814	9	0	0	50	0.0	0.0	1.0	2.0	38.0
2	28	336951	12	0	0	40	1.0	0.0	1.0	2.0	38.0
3	44	160323	10	7688	0	40	1.0	0.0	1.0	2.0	38.0
4	18	103497	10	0	0	30	0.0	1.0	0.0	4.0	38.0

```
In [326]: x3_rf = rf3.fit(x3.drop(['salary'], axis=1), x3.salary)
x3_rf_pred = x3_rf.predict(x3.drop(['salary'], axis=1))
print(classification_report(xt3.salary, x3_rf_pred))
```

	precision	recall	f1-score	support
0.0	0.81	1.00	0.89	12435
1.0	0.98	0.22	0.37	3846
accuracy			0.82	16281
macro avg	0.89	0.61	0.63	16281
weighted avg	0.85	0.82	0.77	16281

```
In [319]: x4 = x3.copy()
x4['occupation'] = enc.fit_transform(df[['occupation']])
x4.head()
```

Out[319]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	salary	Female	Male	marital-status	native-country	occupation
0	39	77516	13	2174	0	40	0.0	0.0	1.0	4.0	39.0	1.0
1	50	83311	13	0	0	13	0.0	0.0	1.0	2.0	39.0	4.0
2	38	215646	9	0	0	40	0.0	0.0	1.0	0.0	39.0	6.0
3	53	234721	7	0	0	40	0.0	0.0	1.0	2.0	39.0	6.0
4	28	338409	13	0	0	40	0.0	1.0	0.0	2.0	5.0	10.0

```
In [327]: xt4 = xt3.copy()
xt4['occupation'] = enc.fit_transform(golden[['occupation']])
xt4.head()
```

Out[327]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	salary	Female	Male	marital-status	native-country	occupation
0	25	226802	7	0	0	40	0.0	0.0	1.0	4.0	38.0	7.0
1	38	89814	9	0	0	50	0.0	0.0	1.0	2.0	38.0	5.0
2	28	336951	12	0	0	40	1.0	0.0	1.0	2.0	38.0	11.0
3	44	160323	10	7688	0	40	1.0	0.0	1.0	2.0	38.0	7.0
4	18	103497	10	0	0	30	0.0	1.0	0.0	4.0	38.0	0.0

```
In [328]: x4_rf = rf3.fit(x4.drop(['salary'], axis=1), x4.salary)
x4_rf_pred = x4_rf.predict(x4.drop(['salary'], axis=1))
print(classification_report(xt4.salary, x4_rf_pred))
```

	precision	recall	f1-score	support
0.0	0.81	1.00	0.89	12435
1.0	0.98	0.22	0.36	3846
accuracy			0.82	16281
macro avg	0.89	0.61	0.63	16281
weighted avg	0.85	0.82	0.77	16281

```
In [320]: x5 = x4.copy()
x5['workclass'] = enc.fit_transform(df[['workclass']])
x5.head()
```

Out[320]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	salary	Female	Male	marital-status	native-country	occupation	workclass
0	39	77516	13	2174	0	40	0.0	0.0	1.0	4.0	39.0	1.0	7.0
1	50	83311	13	0	0	13	0.0	0.0	1.0	2.0	39.0	4.0	6.0
2	38	215646	9	0	0	40	0.0	0.0	1.0	0.0	39.0	6.0	4.0
3	53	234721	7	0	0	40	0.0	0.0	1.0	2.0	39.0	6.0	4.0
4	28	338409	13	0	0	40	0.0	1.0	0.0	2.0	5.0	10.0	4.0

```
In [329]: xt5 = xt4.copy()
xt5['workclass'] = enc.fit_transform(golden[['workclass']])
xt5.head()
```

Out[329]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	salary	Female	Male	marital-status	native-country	occupation	workclass
0	25	226802	7	0	0	40	0.0	0.0	1.0	4.0	38.0	7.0	4.0
1	38	89814	9	0	0	50	0.0	0.0	1.0	2.0	38.0	5.0	4.0
2	28	336951	12	0	0	40	1.0	0.0	1.0	2.0	38.0	11.0	2.0
3	44	160323	10	7688	0	40	1.0	0.0	1.0	2.0	38.0	7.0	4.0
4	18	103497	10	0	0	30	0.0	1.0	0.0	4.0	38.0	0.0	0.0

```
In [330]: x5_rf = rf3.fit(x5.drop(['salary'], axis=1), x5.salary)
x5_rf_pred = x5_rf.predict(xt5.drop(['salary'], axis=1))
print(classification_report(xt5.salary, x5_rf_pred))
```

	precision	recall	f1-score	support
0.0	0.80	1.00	0.89	12435
1.0	0.98	0.21	0.34	3846
accuracy			0.81	16281
macro avg	0.89	0.60	0.62	16281
weighted avg	0.85	0.81	0.76	16281